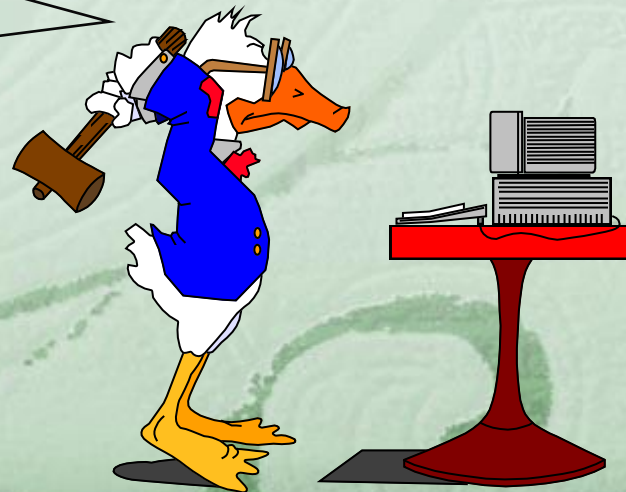


Computer System Overview

Let's figure out
what's inside
this thing...



Multiple Interrupts

■ 2 Choices

☞ Disable Interrupts

- Disable upon entering an interrupt handler
- Enable upon exiting

☞ Allow Interrupts

- Allow an interrupt handler to be interrupted
- Priorities?




Figure 1.12 Transfer of Control with Multiple Interrupts

Multiple Interrupts - Sequential Order

- Disable interrupts so processor can complete task
- Interrupts remain pending until the processor enables interrupts
- After interrupt handler routine completes, the processor checks for additional interrupts
- What happens to the interrupts that occur when disabled?

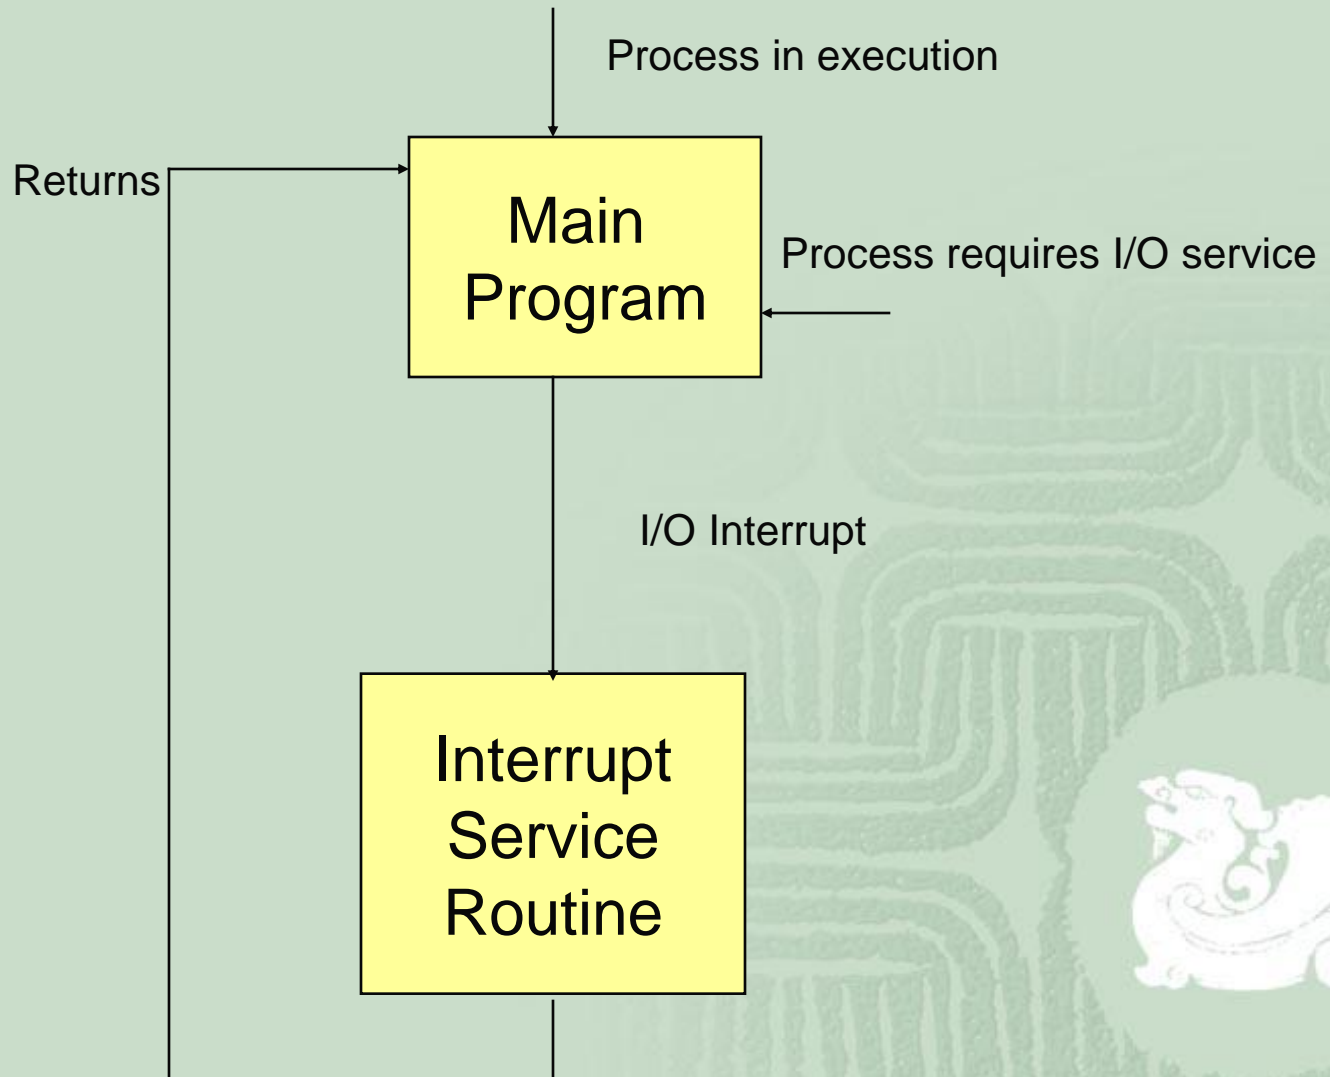


Multiple Interrupts - Priorities

- Higher priority interrupts cause lower-priority interrupts to wait
- Causes a lower-priority interrupt handler to be interrupted
- Example when input arrives from communication line, it needs to be absorbed quickly to make room for more input
- How does this change interrupt handlers?



General flow



Steps in handling interrupts

- Disable further interrupts.
- Store current state of program.
- Execute appropriate interrupt handling routine.
- Restore state of program.
- Enable interrupts.
- Resume program execution.



Examples of interrupts

- Mouse moved/clicked.
- Disk drive operation completed.
- Keyboard key pressed.
- Printer out of paper.
- Video card wants memory access.
- Modem sending or receiving.
- USB scanner has data.



Interrupt handlers

- Interrupt handlers are the routines that are called when an interrupt is detected.
- Interrupt handlers are usually short sections of code that are designed to handle the immediate needs of the device and return control to the operating system or user program.



Interrupt priority

- When multiple I/O devices are present in an interrupt system, two difficulties must be resolved:
 - How to handle interrupt requests from more than one device at a time.
 - Identification of the selected device.
- Assigning priority levels to each device means that highest level priorities will be served before lower levels.



Summary

- An **interrupt** is a signal to the processor emitted by hardware or software indicating an event that needs immediate attention.
- An interrupt can occur at any time.
- Hardware and software are needed to support interrupt handling. The hardware must choose the appropriate time in which to interrupt the executing program and transfers control to an interrupt handler.
- The current state of the interrupted program must be saved.

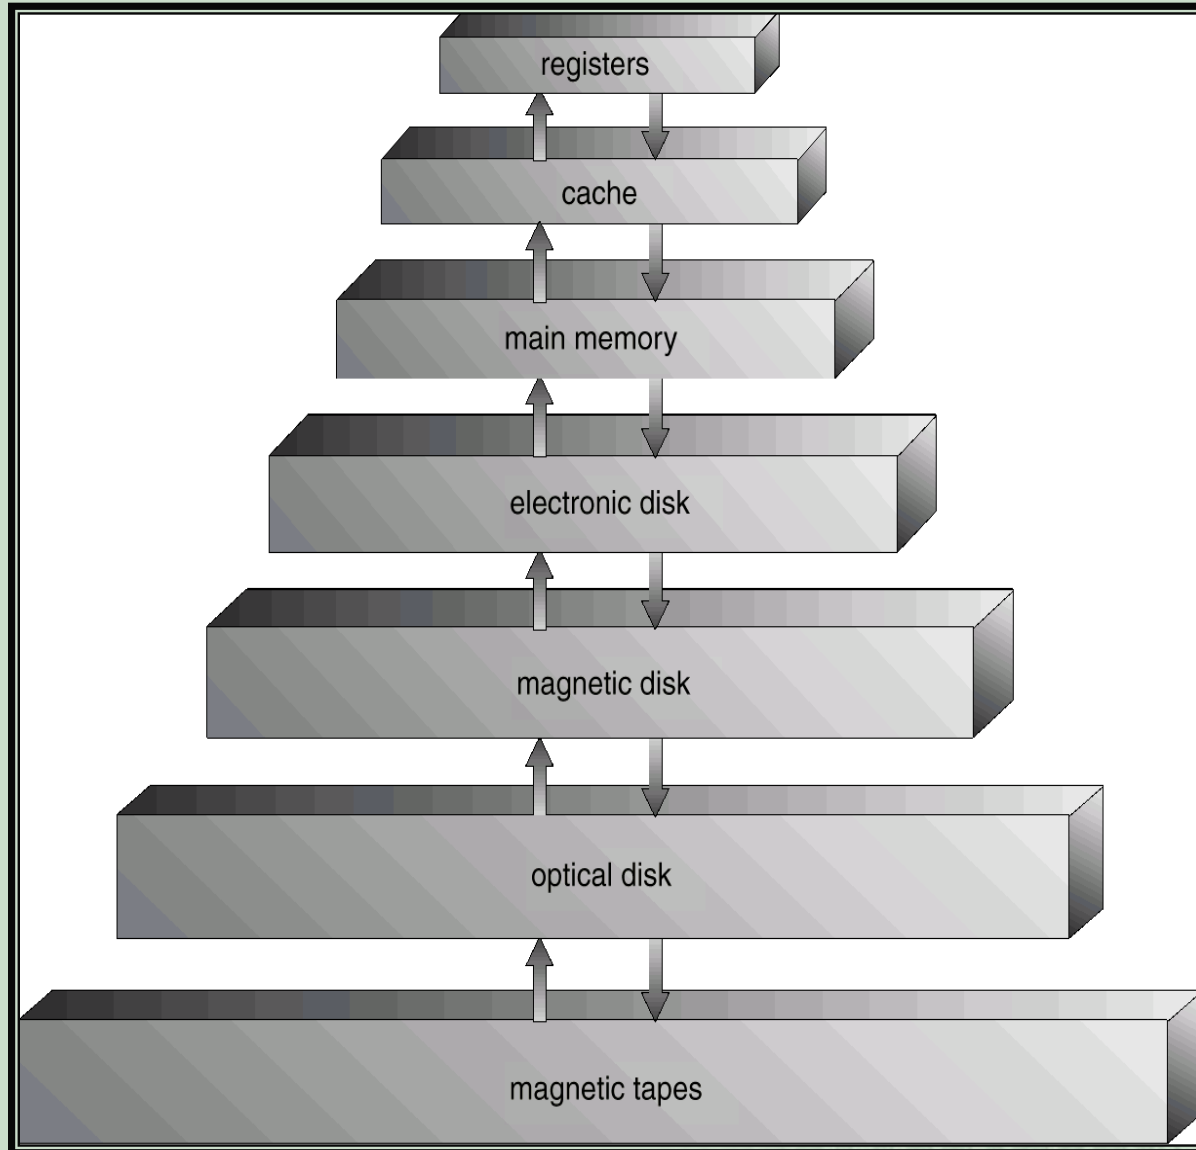


Computer Systems

- Registers
- Interrupts
- **Caching**
- Input/Output
- Protection
- Summary

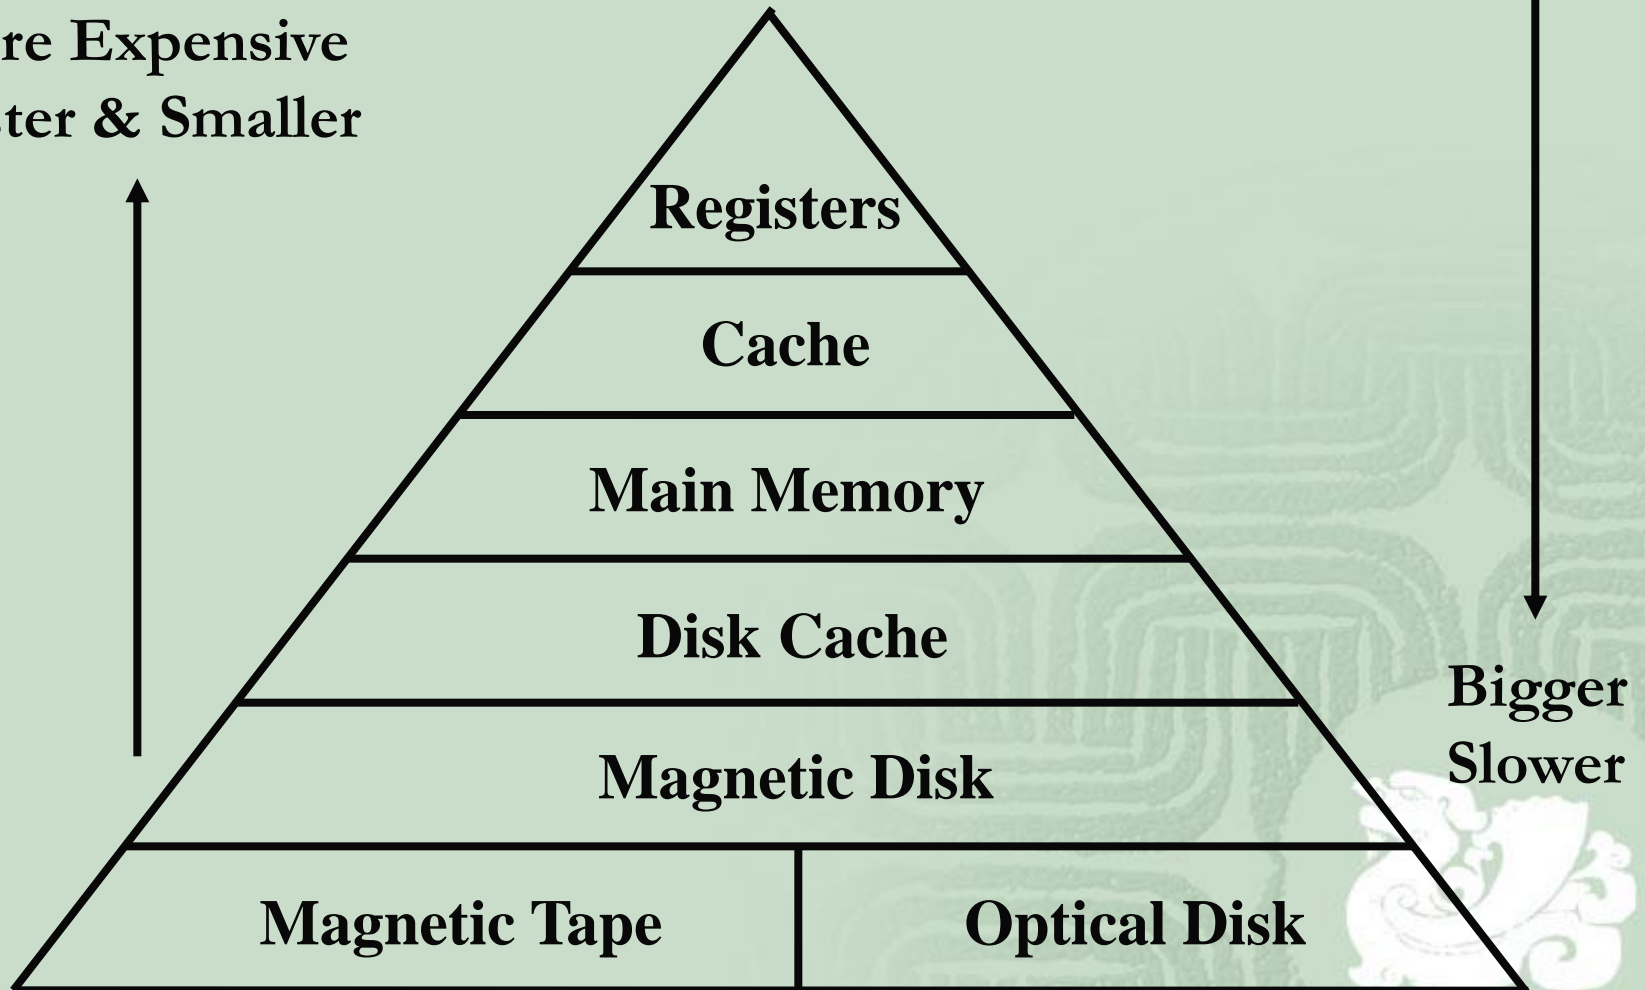


Storage-Device Hierarchy



Memory Hierarchy

More Expensive
Faster & Smaller



Comparative access timings

- Processor registers: 5 nanoseconds
- SRAM memory: 15 nanoseconds
- DRAM memory: 60 nanoseconds
- Magnetic disk: 10 milliseconds
- Optical disk: (even slower)

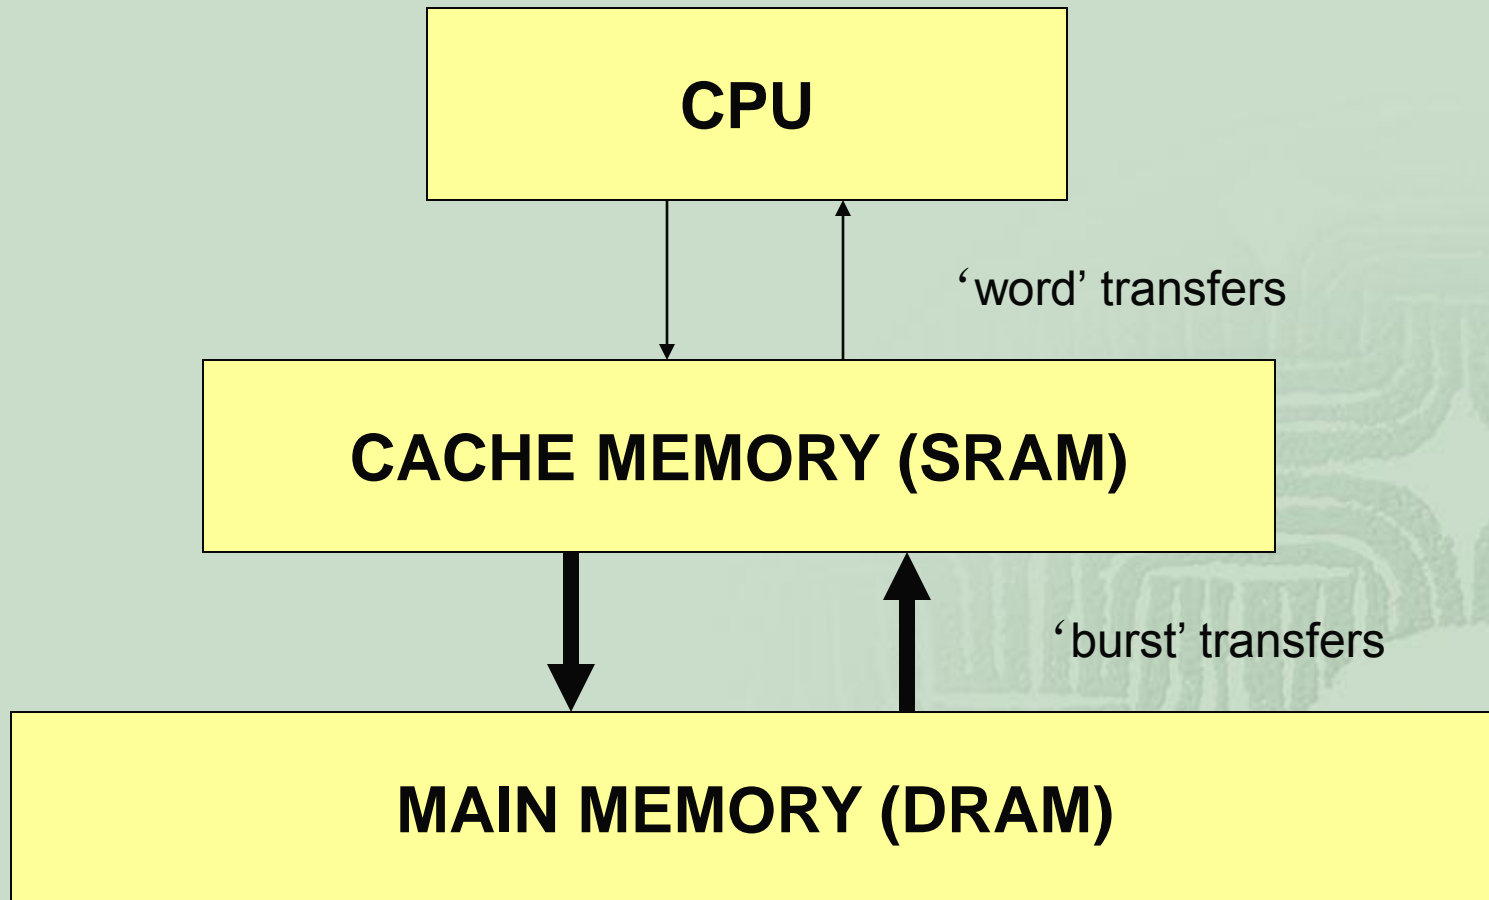


System design decisions

- How much memory?
- How fast?
- How expensive?
- Tradeoffs and compromises
- Modern systems employ a 'hybrid' design in which small, fast, expensive SRAM is supplemented by larger, slower, cheaper DRAM



Memory hierarchy



Memory Cache

- Given:
 - ∞ Processor speed is faster than memory speed
 - ∞ Execution/data localizes
- Cache:
 - ∞ Contains a portion of main memory
 - ∞ Invisible to operating system
 - ∞ Increases the speed of memory
- Processor first checks cache
- If not found in cache, the block of memory containing the needed information is moved to the cache



Cache Design

- Cache size

- ☞ small caches have a significant impact on performance

- Block size

- ☞ the unit of data exchanged between cache and main memory

- ☞ hit means the information was found in the cache



Cache Design

- Mapping function
 - ∞ Determines which cache location the block will occupy
- Replacement algorithm
 - ∞ Determines which block to replace
 - ∞ Least-Recently-Used (LRU) algorithm
- Write policy
 - ∞ write a (dirty) block of cache back to main memory



Another caching technique: Disk Cache

- A portion of main memory used as a buffer to temporarily to hold data for the disk
- Disk writes are clustered
- Some data written out may be referenced again. The data are retrieved rapidly from the software cache instead of slowly from disk



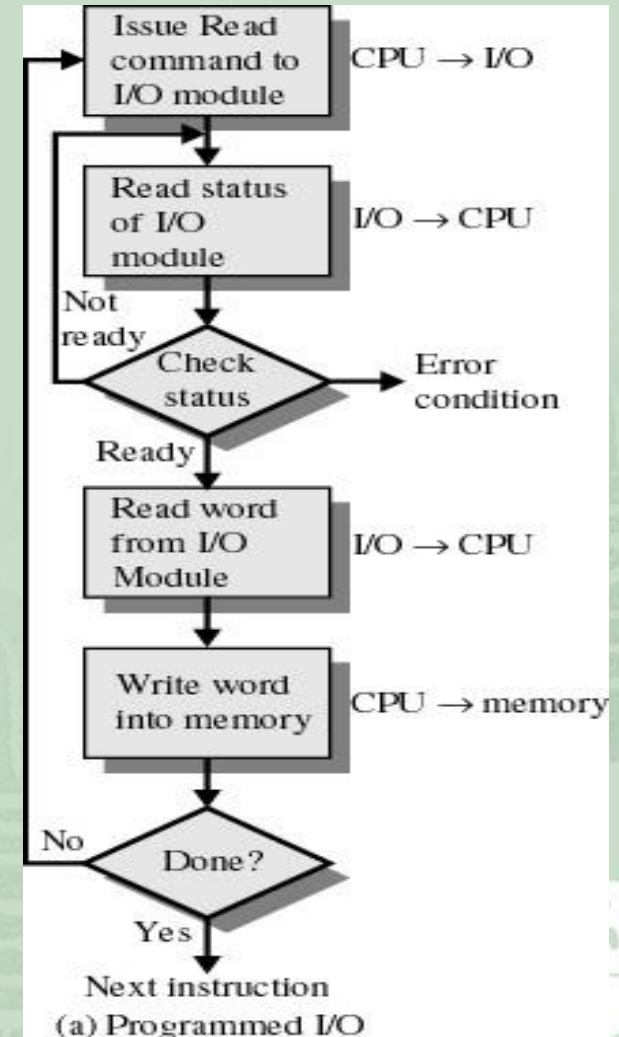
Computer Systems

- Registers
- Interrupts
- Caching
- **Input/Output**
- Protection
- Summary



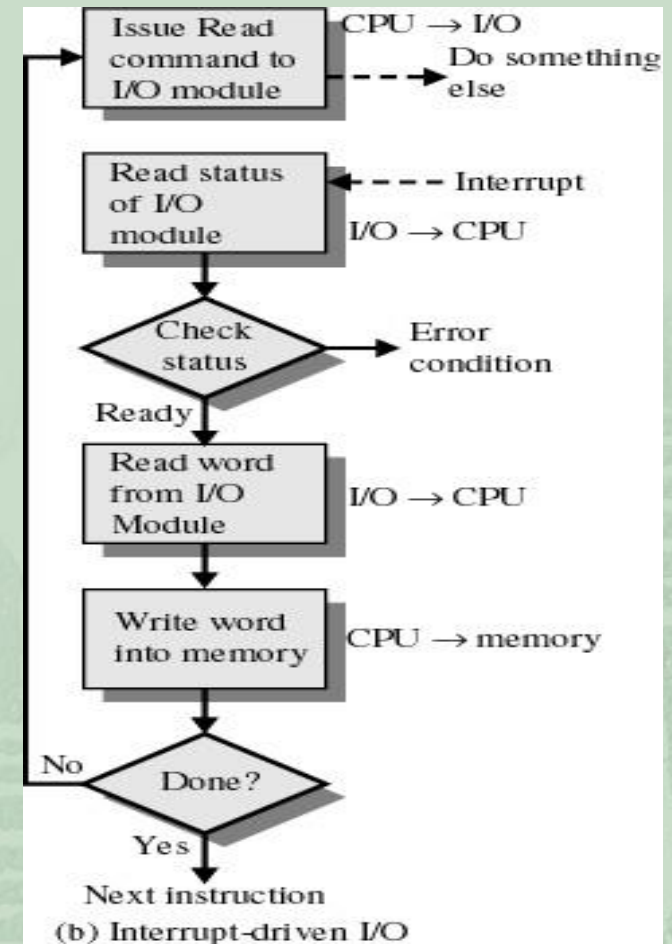
Programmed I/O

- I/O module performs the action, not the processor
- Sets appropriate bits in the I/O status register
- No interrupts occur
- Processor is kept busy checking status (busy waiting)



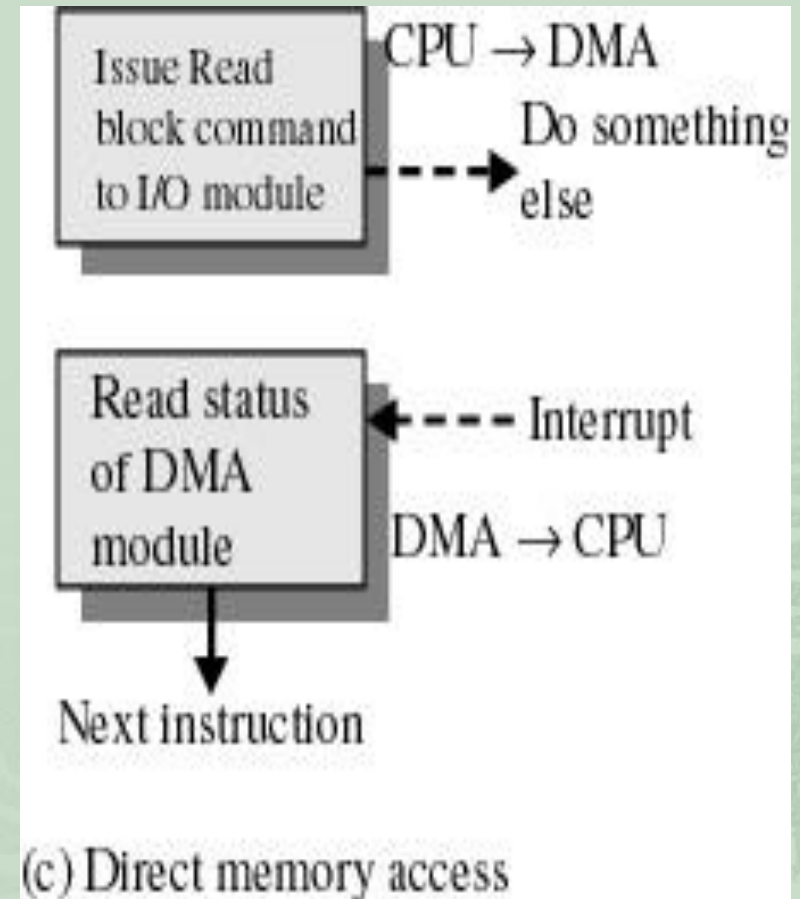
Interrupt-Driven I/O

- Processor is interrupted when I/O module ready to exchange data
- Processor is free to do other work
- No needless waiting
- Consumes a lot of processor time because every word read or written passes through the processor



Direct Memory Access

- Transfers a block of data directly to or from memory
- An interrupt is sent when the task is complete
- The processor is only involved at the beginning and end of the transfer

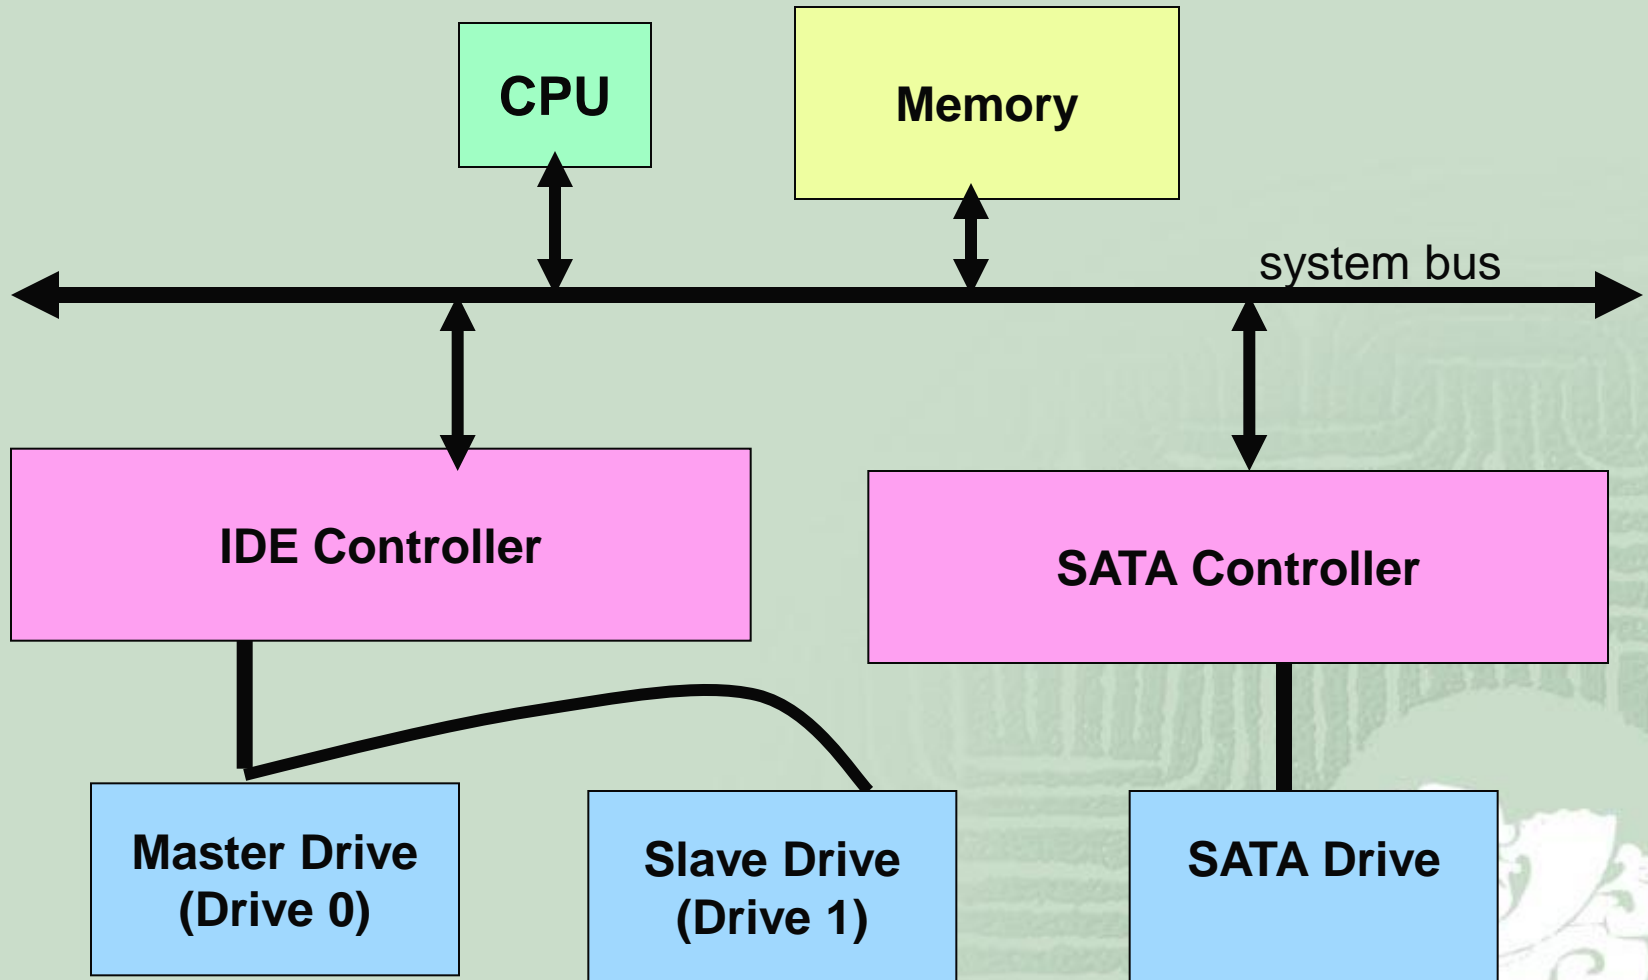


Example: Hard Disks

- Most PCs use either **SATA** or **IDE** disks for storage of files and system data.
- **SATA** stands for **Serial Advanced Technology Attachment** (or **Serial ATA**)
- **IDE** means **Intelligent Drive Electronics**, also called **Parallel ATA** or **PATA**.
- By the beginning of 2007, **SATA** had largely replaced **IDE** in all new systems.
 - Inexpensive, large storage capacity, hot plugging, faster



The HD Controller



‘IDENTIFY DRIVE’

- There exist about 40 different commands (e.e., read, write, seek, format, sleep, etc)
- Some are ‘mandatory’, others ‘optional’
- An example: the ‘Identify Drive’ command
 1. It provides information on disk’s geometry and some other operational characteristics
 2. It identifies the disk’s manufacturer and it provides a unique disk serial-number



Disk Command Protocol

- Disk Commands typically have 3 phases:
 - ⌘ COMMAND PHASE: CPU issues a command
 - ⌘ DATA PHASE: data moves to/from buffer
 - ⌘ RESULT PHASE: CPU reads status/errors



Computer Systems

- Registers
- Interrupts
- Caching
- Input/Output
- **Protection**




Hardware Protection

- Why protect hardware?
- From what?
 - ∞ Shared hardware resources – memory, disk, ...
 - ∞ Errant programs
- How?
 - ∞ CPU provides 2 modes of operation
 - User Mode (non-privileged)
 - Supervisor/monitor/OS mode (privileged)
 - ∞ Privileged instructions can only be executed in monitor mode
 - All I/O instructions are privileged




Dual-Mode Operation (Cont.)

- *Mode bit* added to computer hardware to indicate the current mode: monitor (0) or user (1).
- When an interrupt or fault occurs hardware switches to monitor mode.

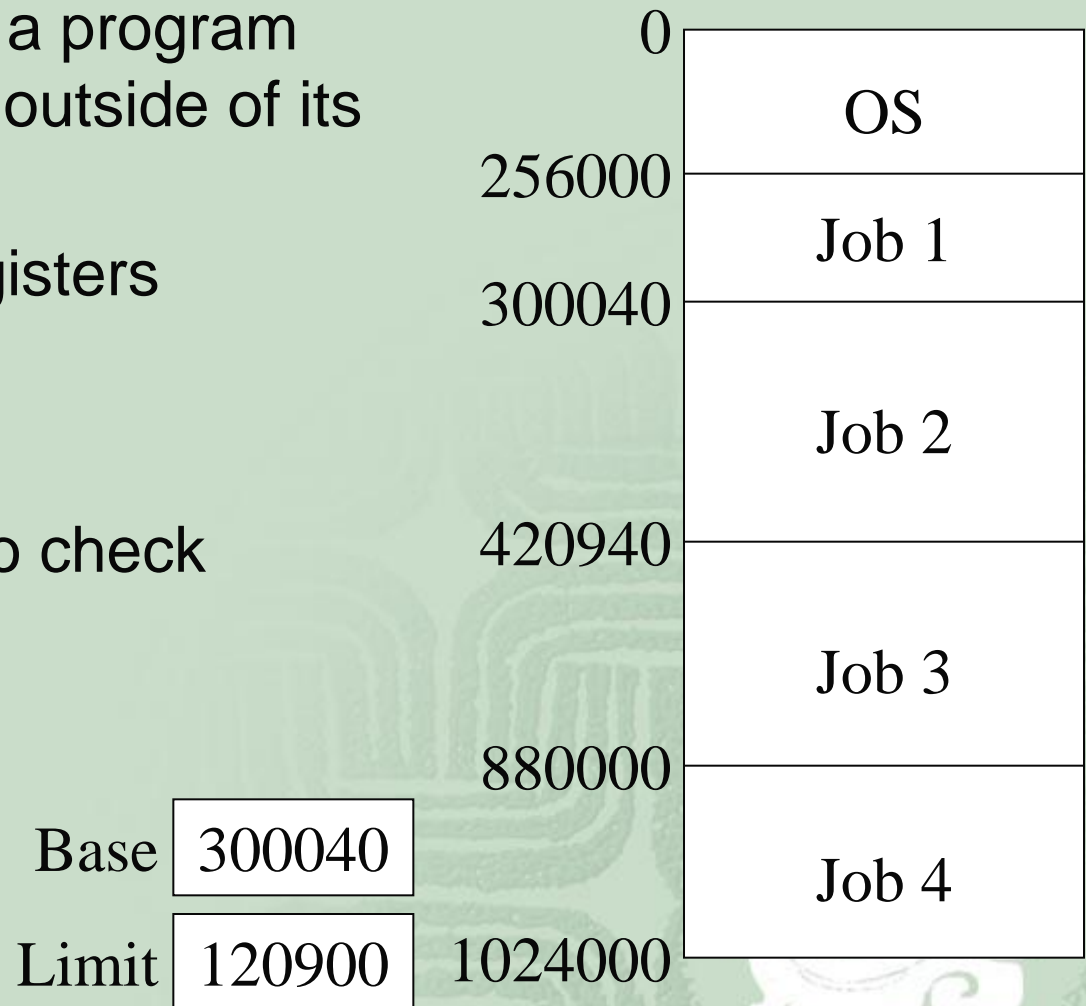


Privileged instructions can be issued only in monitor mode.




Memory Protection

- We must not allow a program access to memory outside of its space.
- How? Add two registers
 - ↻ Base
 - ↻ Limit
- Use the registers to check every reference



Hardware Address Protection



CPU Protection

- What is there to protect?
 - ❧ Configuration
 - ❧ CPU as a resource
- Timer
 - ❧ Task Switching uses the timer
 - When process is loaded, timer is set
 - Timer is decremented each cycle
 - When time is zero, an interrupt is generated
 - Process is switched
 - ❧ Load-timer is a privileged instruction



I/O Protection

- All I/O instructions are privileged instructions.
- Must ensure that a user program could never gain control of the computer in monitor mode.



Protect at all cost?

- If all I/O instructions are protected, how do you perform input or output?
- System calls
 - ✧ Often the call is a trap (soft interrupt) to a ISR (Interrupt Service Routine)
 - ✧ Control is passed to the ISR which sets the mode bit to **monitor mode**
 - ✧ ISR verifies that parameters are correct
 - ✧ ISR executes request, resets mode
 - ✧ Control is returned to user program

